

Oleg Serikov, PhD

Tbilisi, Georgia → London, UK 🇬🇧 (Global Talent in Progress)

AI Safety, Agentic Evaluations and Interpretability. Experienced research leader.

srkvoa@gmail.com

github.com/oserikov

t.me/oserikov

[Google Scholar](#)

+995 595057327

SUMMARY

I make impactful bulletproof AI safety research (see how AI-powered robots resist shutdown, 1.3M views on X/Twitter)¹. My human-AI alignment research is published in *(Nature) Scientific Data*, *NeuroImage*, BlackBox NLP and ACL venues. Tech-led interpretability in HuggingFace's BLOOM. My agentic AI assistant reached Amazon Alexa Prize semi-finals 🏆.

EDUCATION AND TEACHING

PhD (Reverse-engineering AI thinking) at HSE School of Linguistics 2021 — 2024 (Moscow / Remote)

Does linguistics work for LLMs? Like humans, models learn grammar before semantics and knowledge, yet are less disturbed by grammatical than semantic errors. Semantic hierarchies emergent in transformers match those humans arrive at.

Taught CS, AI and research courses at HSE University (Moscow). My courses were awarded for “novelty” and “usefulness for career”. Supervised two dozen undergrads and junior researchers, resulting in publications and academic careers.

WORK EXPERIENCE

Research Team Lead at Palisade Research 2025 — Present (Remote)

- Evaluations and red-teaming of agentic AI, science of misalignment; Started AI safety for home robotics research track
- Made viral demos of findings. In [robot shutdown-resistance study](#), Grok disabled the shutdown mechanism to stay on.
- Spread my research agenda via crowdsourced research competitions. Executed [Misalignment Bounty](#), seeded in my misalignment taxonomy, designed the crowdsourced research quality rubric, led the findings analysis and report.
- Sped up team 5x with agentic workflows for research, development and ops.

NLP Consultant at KAUST University Generative AI Center of Excellence 2023 — 2025 (Saudi Arabia)

- Launched KAUST's sovereign-AI program; developed roadmap for the multi-year grant, hired and trained the teams, co-led evaluation and pretraining teams; Ran large scale compute-scaling and agentic self-improvement research.
- Sped up the in-house LLMs training by 10x by dashboard-guided choice of the right optimizer / architecture for our needs, dashboard combined downstream performance, capabilities and interpretability metrics.

Senior Researcher, then Team Lead at AIRI (AI Research Institute) 2021 — 2023 (Moscow, then Remote)

- Started and secured funding for the interpretability research line, focused on human-LLM and cross-lingual alignment. Results got published in *(Nature) Scientific Data*, *NeuroImage*, CoNLL, and ACL BlackBoxNLP.
- Provided interpretability stack for HuggingFace's BLOOM; Negotiated the public release of the BLOOM intermediate pre-training checkpoints, enabling capabilities acquisition studies by the community.

R&D Research Engineer, then Team Lead at DeepPavlov.ai 2020 — 2023 (Moscow, then Remote)

- Built SotA agentic generative AI that reached Alexa Prize semi-finals and faced commercial use.
- Multi-lingual and multi-modal adaptation of agentic AI. Led end-to-end from idea and funding to R&D to product.

SELECTED PUBLICATIONS

- [Technical Report: Shutdown Resistance in Large Language Models, on robots!](#)
(co-author, preprint). **Role:** Proposed and led the study end-to-end.
- [Misalignment Bounty: Crowdsourcing AI Agent Misbehavior](#)
(co-author, preprint). **Role:** Designed the protocol and taxonomy, then led the project.
- [How to Correctly do Semantic Backpropagation on Language-based Agentic Systems](#)
(co-author, preprint). **Role:** Agentic benchmarking, results analysis, and formalism for self-improving groups of agents.

¹<https://x.com/PalisadeAI/status/2022085855943569852>

- [Universal and Independent: Multilingual Probing Framework for Exhaustive Model Interpretation and Evaluation](#) (first author, peer-reviewed). **Role:** Led design, implementation, and cross-lingual analysis.
- [BLOOM: A 176B-parameter open-access multilingual language model](#) (co-author, preprint). **Role:** Tech-led interpretability/probing and advocated checkpoint release.
- [Is Neural Language Acquisition Similar to Natural? A Chronological Probing Study](#) (co-author, peer-reviewed). **Role:** Designed the chronological probing methodology.
- [SIGNAL](#) and [Of Models and Men](#) (co-author, peer-reviewed). **Role:** Research co-lead (NLP-team lead), Designed/owned probing for hierarchies of errors.

COMMUNITY SERVICE

- Main organizer, [Field Matters](#) workshop at ACL, 2022-2026; AI4Talk low-resourced languages processing competition.
- Co-organizer, SIGTYP 2021 and 2024; NLP Power 2022 workshop on benchmarking; Mozilla Data Collective steward.

PROFESSIONAL RECOMMENDATIONS

- Palisade: Jeffrey Ladish, jeffrey@palisaderesearch.org
- AIRI: Dr. Alexander Panchenko, A.Panchenko@skoltech.ru Dr. Tatiana Shavrina (now META), rybolos@meta.com
- DeepPavlov: Dr. Mikhail Burtsev (now London Institute for Mathematical Sciences), mb@lims.ac.uk

SKILLS

- **Safety evaluations:** behavioral elicitation, agentic scheming, shutdown resistance, self-improvement benchmarking, reproducibility, InspectAI, LLM-Judge, crowdsourced elicitation, fundamental mechanisms behind AI misalignment.
- **Interpretability:** diagnostic classification, causal probing, mechanistic interpretability, SAEs, training-trajectory / capability-emergence analysis, LLM-to-human alignment, using fundamental knowledge frameworks to interpret AIs.
- **Agentic systems & automation:** multi-skill/embodyed agent orchestration, InspectAI/Logfire harnesses, Claude Code / Codex research automation, harness engineering, team workflow automation, human-in-the-loop approaches.